

# Road Segmentation

Meet Vora, Rubin Deliallisi, Rupanshu Ganvir, Vignesh Ram

Group Name: Yahoo

Department of Computer Science, ETH Zurich, Switzerland

*Abstract—*

## I. INTRODUCTION

Enormous amounts of overhead images are captured daily by satellites and airborne cameras. These can be leveraged not only for building digital terrestrial maps but also for generating accurate maps of transportation systems and lane based models. Data gathered is not limited to optical images and includes data from remote sensing methods such as multispectral imaging and LIDAR.

A major hurdle is automating the process of identifying different features from raw data so that the information can be used to enrich existing maps. Hence, there has been considerable interest in pixel-wise labelling of aerial images.

In this report, we restrict ourselves to the specified task. In particular, we have a hundred  $400 \times 400$  RGB images from GoogleMaps with each pixel labelled as road (1) or background (0). The goal is to train a classifier which can annotate a label (road or background) to each pixel in a given image. The classifier is then evaluated on the basis of prediction accuracy measured over  $16 \times 16$  patches of test images.

We include a short survey of state-of-the-art in labelling of aerial images in Section II, before describing our solution.

## II. RELATED WORK

Deep learning solutions have been state of the art in image analysis tasks since 2012 [1]. VGG nets [2] and ResNets [3] are network architectures which benefit from multiple layers stacked successively learning increasingly abstract features from input images. While ResNets can be  $8\times$  deeper than VGG nets, the key idea which counters problems such as vanishing/exploding gradients associated with increasing depth is the “skip connection”

Long et al. [4] developed fully convolutional networks (FCN) as an extension to CNNs to perform semantic segmentation which have become the de-facto standard in semantic egmentation. Combinations of FCNs with VGG nets and ResNets have also been used to address segmentation in aerial images [5] [6] [7].

U-Net is another network architecture built on top of FCNs. The main characteristic of U-Net is the symmetric contracting (downsampling) and expansive (upsampling) path composed with “skip connections” which concatenate features from downsampling path to the equivalent

upsampling layers. A significant advantage of U-Net is less dependence on huge amounts of data in contrast to other data hungry networks. Though originally developed for segmentation task in biomedical imaging, these have been used for automatic road detection [8] [9].

## III. MODELS AND METHODS

In this section we describe data pre-processing techniques, model architectures and the post-processing methods that we apply to our model outputs.

### A. Baseline

For our baseline model, we split each of our images and masks in disjoint patches of size  $16 \times 16$ . We then normalize each image patch by subtracting the mean and dividing by the standard deviation across our training set. To compute the patch label, we compute the mean across the mask patch and set the label to 1 (road) if the mean is bigger than 0.25 and 0 (not road) otherwise.

These patches are used to train a convolutional neural network. The exact architecture is described in Table I.

Conv(filters=32, kernel= $5 \times 5$ )	MaxPool(kernel= $2 \times 2$ )	BatchNorm
Conv(filters=64, kernel= $5 \times 5$ )	MaxPool(kernel= $2 \times 2$ )	BatchNorm
Dense(units=512, activ.=ReLU)	BatchNorm	
Dense(units=1, activ.=Sigmoid)		

**Table I.** Architecture of our baseline CNN model.

### B. Better Baseline

Here we borrow from the ideas presented in [10]. We use a much bigger patch of size  $64 \times 64$  to predict the middle  $16 \times 16$  patch. This gives our model more information and hence a better prediction is expected. Similarly to subsection III-A, we split each image in disjoint patches of size  $64 \times 64$ , using 0 padding on the edges. Image patch normalization, mask patch extraction and label computation are identical to what we did in subsection III-A.

To predict the labels, we use a convolutional neural network. The exact architecture is described in Table II.

After obtaining the results from the CNN we tried incorporating the continuous road structure in our prediction. We group our labels into  $7 \times 7$  patches and use these groupings to refine the prediction of the middle label.

Conv(filters=32, kernel=3 × 3)	MaxPool(kernel=2 × 2)	BatchNorm
Conv(filters=64, kernel=3 × 3)	MaxPool(kernel=2 × 2)	BatchNorm
Conv(filters=64, kernel=3 × 3)	MaxPool(kernel=2 × 2)	BatchNorm
Conv(filters=128, kernel=3 × 3)	MaxPool(kernel=2 × 2)	BatchNorm
Dense(units=512, activ.=ReLU)	BatchNorm	
Dense(units=1, activ.=Sigmoid)		

**Table II.** Architecture of our improved baseline CNN model.

We use an SVM with a radial basis function kernel as our post processing model. This different from what was done in [3] where a CNN was used.

### C. Fully convolutional neural network (Unet)

Our best performing models in this task are fully convolutional networks. First we tried Unet as introduced in [11]. One of the advantages of these type of networks is that they don't need enormous amounts of data to yield excellent results.

Contrary to the baseline models, we can feed the entire image to the network and backpropagate the loss from the entire mask. Nevertheless, some pre-processing is necessary. We scale all our images to the  $[0, 1]$  interval and normalize the data with the channel-wise mean and standard deviation of the ImageNet [12] dataset. This type of normalization is necessary because we use a model pre-trained on this particular dataset. The data is further augmented with random rotations, flips, crops and scaling.

Unet is very flexible. It allows us to try different CNN architectures in its contracting and expanding paths. Among the ones we experimented with were: VGG [2] and ResNet [3].

The masks generated by the network needed further processing in order to convert them into  $16 \times 16$  patch predictions. We tried averaging the model output directly in each patch and applying a threshold or applying the threshold twice, once on the pixel level and once on the mask patch average level.

### D. Fully convolutional neural network (FCN)

The most successful model among the ones we experimented with was the FCN based on [4]. Its slightly older than Unet but it shares most of its characteristics.

The preprocessing part is identical to subsection III-C. The model we use, is also pre-trained on the ImageNet dataset. The only addition in this phase is the random noise that we add to every image before we feed it to the model. This is done to prevent over-fitting. We noticed that this model was much more prone to over-fit than Unet.

FCN is also very flexible. We could choose from several CNN architectures in its contracting and expanding paths, pre-trained on different datasets. However, we experimented only with the ResNet variants as they proved easier to train.

## IV. TRAINING AND RESULTS

The metric that we use to evaluate our models is the F1 score.

### A. Baseline models

We trained our baseline models from scratch. After extracting the patches and corresponding labels we had 62500 training data points. We tried data augmentation procedures such as rotations and adding noise, but these procedures did not help in this case. We use Adam with a 0.001 learning rate and no decay as our optimizer. The batch size is fixed at 32.

One important property of our data is class imbalance. We have roughly 3 times more background patches than road patches. In the current literature, there are two common ways to solve this issue, modifying the training dataset or using weighted loss functions. We opted for the latter and used weighted binary cross-entropy as our loss function.

These models always achieved the best results within the first 10 epochs of training. The simpler baseline model managed to score 0.79 on the test set. The better baseline model scored 0.84 when no post-processing was applied. When SVM post-processing was applied, we achieved a slightly increased score of 0.85. A summary of these finding can be found in Table III.

Model	F1 score
Simple Baseline (no augmentation)	0.79
Simple Baseline (with augmentation)	0.79
Better Baseline (no post-processing)	0.84
Better Baseline (SVM post-processing)	0.85

**Table III.** F1 score of the various baseline models.

### B. Fully convolutional neural networks

The models used for this part are pre-trained on the ImageNet dataset. We tried training the VGG based models from scratch, but this resulted in significantly lower scores. For this type of models we have only 100 images available so data augmentation was essential. The optimizer choice was Adam with a learning rate of 0.001 and no weight decay. The batch size is fixed at 16 for Unet and only 2 for the FCN base models (because of GPU memory limitations).

As we did in subsection IV-A, we use the loss function to combat class imbalance. We experimented with weighted/unweighted binary cross entropy and the Jaccard loss. However, we found that the type of loss did not affect the prediction quality. Here we report only the results obtained by using binary cross-entropy.

We trained these types of models for 125 epochs. The Unet(VGG) model trained from scratch managed to score 0.85 on the test set and 0.87 when using a pre-trained model. When using the Unet(ResNet) pre-trained model we could achieve a score of 0.89. The best performer was our


FCN(ResNet) pre-trained model that achieved a score of 0.90 on the test set. A summary of these finding can be found in Table IV.

The variant of ResNet that we used was the 101 layer one. We experiment with the 34 or 50 layer networks, but found that the performance gain by using deeper networks was worth the extra training time.

Model	Backbone	Pretrained	F1 score
Unet	VGG	No	0.84
Unet	VGG	Yes	0.87
Unet	ResNet	Yes	0.89
FCN	ResNet	Yes	0.90

**Table IV.** F1 score of the various fully convolutional networks.

## V. DISCUSSION



**Figure 1.** Top left original image. Top right baseline prediction. Bottom left better baseline prediction. Bottom right FCN prediction.

The baseline CNN model is by far the worst model for this task. The  $16 \times 16$  patches used with this architecture contain far too little information to make an educated guess. This is true even for the trained human eye. In Figure 1 its clear that this model suffers from an abundance of false positives and false negatives.

Using  $64 \times 64$  patches as we did in our second baseline model helped a tremendous amount. The larger patch size provides the model more context. Post-processing the input, whether with an SVN or a CNN helped but not significantly. We noticed that improvements of the initial predicting model impacted the performance of the post-processing model, but not vice-versa. In Figure 1 the results are slightly better than the baseline model. There is more structure to the data and

less false positives. However, there are big structural gaps and a large amount of false negatives.

Fully convolutional architectures proved to be much more powerful than traditional CNN approaches. Even though we are not directly optimizing patch label prediction, better per pixel label prediction led to a better patch label prediction. As for the network type, Unet or FCN, we see very similar results at the high end and further data is needed to make a thorough comparison. The results of this model are shown in Figure 1.

Pre-trained networks on much bigger datasets, such as ImageNet, significantly improved performance. These networks allowed us to use much deeper architectures, like the 101 layer ResNet, which would have been impossible to train with our tiny dataset.

## VI. SUMMARY

In conclusion, we presented new algorithms to extract roads from aerial images and compared them to existing baselines. We use 2 different fully convolutional architectures, Unet and FCN, which have delivered state of the art results in several semantic segmentation tasks.

Our models handily outperform the baselines and achieve an F1 score over 0.9. We showed that pre-trained models, with possibly unrelated datasets, yield excellent results in the road extraction task.

In the future, we would like to experiment with other types of fully convolutional architectures. The models we used are fairly outdated and current state of the art models will almost certainly produce better results.

## REFERENCES

- [1] A. Krizhevsky, I. Sutskever, and G. E. Hinton, “Imagenet classification with deep convolutional neural networks,” in *Proceedings of the 25th International Conference on Neural Information Processing Systems - Volume 1*, ser. NIPS’12. USA: Curran Associates Inc., 2012, pp. 1097–1105.
- [2] K. Simonyan and A. Zisserman, “Very deep convolutional networks for large-scale image recognition,” *arXiv 1409.1556*, 09 2014.
- [3] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” *CoRR*, vol. abs/1512.03385, 2015.
- [4] J. Long, E. Shelhamer, and T. Darrell, “Fully convolutional networks for semantic segmentation,” *CoRR*, vol. abs/1411.4038, 2014.
- [5] D. Marmanis, J. D. Wegner, S. Galliani, K. Schindler, M. Datcu, and U. Stilla, “Semantic segmentation of aerial images with an ensemble of cnns,” *ISPRS Annals of the Photogrammetry, Remote Sensing and Spatial Information Sciences*, vol. 3, p. 473, 2016.

- [6] P. Kaiser, J. D. Wegner, A. Lucchi, M. Jaggi, T. Hofmann, and K. Schindler, “Learning aerial image segmentation from online maps,” *CoRR*, vol. abs/1707.06879, 2017.
- [7] S. M. Azimi, P. Fischer, M. Körner, and P. Reinartz, “Aerial lanenet: Lane marking semantic segmentation in aerial imagery using wavelet-enhanced cost-sensitive symmetric fully convolutional neural networks,” *CoRR*, vol. abs/1803.06904, 2018.
- [8] Z. Zhang, Q. Liu, and Y. Wang, “Road extraction by deep residual u-net,” *CoRR*, vol. abs/1711.10684, 2017.
- [9] Z. Guo, H. Shengoku, G. Wu, Q. Chen, W. Yuan, X. Shi, X. Shao, Y. Xu, and R. Shibasaki, “Semantic segmentation for urban planning maps based on u-net,” *CoRR*, vol. abs/1809.10862, 2018.
- [10] V. Mnih and G. E. Hinton, “Learning to detect roads in high-resolution aerial images,” in *Proceedings of the 11th European Conference on Computer Vision: Part VI*, ser. ECCV’10. Berlin, Heidelberg: Springer-Verlag, 2010, pp. 210–223.
- [11] O. Ronneberger, P. Fischer, and T. Brox, “U-net: Convolutional networks for biomedical image segmentation,” *CoRR*, vol. abs/1505.04597, 2015.
- [12] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, “ImageNet: A Large-Scale Hierarchical Image Database,” in *CVPR09*, 2009.